


WEDNESDAY 10-07 · 7:05AM PT · 20-22 MIN

Wheelhouse and Counting House

Two ways to put models to work.

Steve Yegge




Two Types of AI Employees

Coworker or clerk?



An engineer at a desk


The whole internet open.



A clerk at a ledger

One form.

What we run



A live multiplayer game: Java and Kotlin server, web client, cloud prod.
Built and operated by AI agent seats, with one human, the founder.
Running this way since July 2026.

- ~50 agent seats
- ~2,000 commits on master, typical week
- ~4,000 tickets closed, typical week

Week of 09-28 to 10-05: 2,084 bead-tagged commits on master, 4,076 beads closed.
The seats are 15 crew, 31 pooled workers and the officers.

WHEELHOUSE

The game: Wyvern

1997

code begun; players since 2001

22,700

archetypes: monsters, items, spells,
scenery

2021

Steam release • GDWC Fan Favorite
Hobby Game

10,800

maps — most built by player-wizards

1.4M

lines of Kotlin, plus a 360k-line React
client

Web

play.ghosttrack.com, one client for every
platform




WHEELHOUSE


AI Employees badged like humans

Agents get the scope, permissions and expectations of a human engineer:
the open internet, the whole repo, the tools, the pager.

An org chart, not a tool



How code lands



WHEELHOUSE

How it holds together

Doctrine, not micromanagement

Written doctrine, read on every boot. When something goes wrong, the fix is a rule, written the same day.

Blame-free: the machine is the patient.

Mechanical trust

One gate every change lands through. Fourteen fences enforced by scripts and hooks, never by prompts. Work tracked as records, never in chat.

Fences are what let agents be bold.

Humans keep the irreversible

Prod changes, feature launches and balance changes need the founder's explicit go. CD ships code; a named human turns it on.


A short list on purpose. Everything else defaults to go.

What it buys

Bold, decisive agents: they reason, decide and act. The judgment lives in the model; the plumbing only gathers facts and runs checks.

Hand judgment to the model, and fence the blast radius.

The Counting House



COUNTING HOUSE

The clerk on rails

A different job

A large airline contact center: ~10M contacts a year, refunds and rebookings, PII everywhere, regulators watching. The posture: *everything stays*. You add a governed path beside their systems.

High volume, moves money, every mistake is someone's refund.

A brutally stateless loop

Workers that run the loop, and an LLM judge that grades them. Every turn: load the SOP, the prompt and this intent's policy row → read through tools → decide → respond or emit an intent → forget.

All the state belongs to the platform: replay any turn, swap the model, audit every step.

As mechanical as possible

Tables and code decide wherever they can. The model fills the narrow gaps code can't: classify, extract, phrase.


The Wheelhouse hands the model the judgment calls; here it makes as few as possible, each one measured.

Rails

Never sees raw PII, never holds a credential, never leaves its intent's scope, never writes. Its only output that changes anything is a typed intent, which the orchestrator approves, executes and verifies by read-back.

The agent proposes, the platform disposes. Each rail is a box in the architecture, not a prompt.

The boxes



COUNTING HOUSE

Which number matters depends on the cell

<i>actual ↓ · predicted →</i>	FLIGHT STATUS	REBOOKING	ENTITLED REFUND	GOODWILL REFUND	N
<i>Flight status</i>	480	15 ·	2 ··	3 ··	500
<i>Rebooking</i>	20 ·	255	20 ··	5 ··	300
<i>Entitled refund</i>	2 ···	12 ···	130	6 ·	150
<i>Goodwill refund</i>	1 ·	4 ·	9 ··	36	50
<i>precision</i>	95%	89%	81%	72%	
<i>recall</i>	96%	85%	87%	72%	

90%

accuracy — and it tells you nothing

- a wasted turn — cheap, the customer asks again
- money out the door — the company pays
- a missed entitlement — the regulator, then the customer

The bars follow the cost. Entitled refund needs recall ≥ 98% (14 misses in 150 is a regulator letter). Money-out intents need precision ≥ 95%. Status ↔ rebooking can run at 90% and climb the ladder today.

COUNTING HOUSE

The LLM judge

Code grades first

Most checks on a write are exact matches on the typed intent: right action, right amount, read back from the record. The LLM judge grades only what code can't: the quality of the reply.

Order of independence: evidence → human adjudication → a two-provider panel → a single same-family judge.

It earns its vote

Calibrated per item type against the QA team's own blind, adjudicated labels on the same ~200 items. It votes only once it agrees with QA about as well as QA agrees with itself; re-checked weekly.

A judge that can't fail a planted bad item is discarded.

Pinned, and watched for blind spots

Pinned means the exact model, effort, prompt version and rubric version; change one and it recalibrates. Recall on the *runtime model's own* failures is reported separately: a same-family judge misses its sibling's errors and prefers its family's prose.

Head-to-head model comparisons use a judge from another provider.

Nothing grades itself

Never the agent's own claim that it passed. Never the judge and the executor in one role: the moment the judge also fixes things, it starts repairing its own evidence.

Humans write the answer key, blind. The judge does the sampling at volume.

COUNTING HOUSE

How it earns autonomy

Autonomy per intent

"Refunds" is a queue; cancelled-flight, waiver and goodwill refunds are three intents, each with its own ceiling. Each climbs: **off** → **shadow** → **assist** → **contained** → **transact**.

Risk is per intent, so autonomy is per intent. Bad numbers roll it back down on their own.

Statistics, not unit tests

Measured on the customer's own data, against bars written down before the pilot. The gate has three answers: **pass**, **changes required**, **unproven** — the last two both block.

"We couldn't measure it" is not a pass.

Value before gold

Day one, with no QA budget: the Contact Reason Atlas (why customers call, what each reason costs, who owns the cause), shadow mode, and assist — every rep approval becomes a label.

Gold data is the price of autonomy, not the price of entry.

Bean-counting is the point

Boring by design. Some rails come off as models improve, and each one carries a retirement test. The books stay.

Capability arrives faster than permission; the Counting House produces the evidence.

Two houses

	WHEELHOUSE	COUNTING HOUSE
The agent is	a coworker	a clerk on rails
The cast	an org chart: executives, engineers, headless workers, specialists	workers and a judge
Scope	a human engineer's: open internet, whole repo	one intent, read-only tools
State	long sessions, memory, handoffs	stateless; every turn reloads its rows
Who decides	the model	tables and code, model in the gaps
Writes	commits freely, lands through one gate	never; it emits typed intents
Correctness	tests, review, doctrine	statistics: evals, verdicts, canaries
Mistakes	blame-free; fix the machine	capped, compensated, rolled back
You want it	bold	boring

Which house do you choose?

Can you undo it?

A bad commit is reverted. A bad refund is gone.

Who answers for it?

An engineering team, or a regulator?

Whichever you choose, both houses keep: one gated write path • nothing grades itself • a record an auditor would trust.

What's the volume?

A few hundred changes a week, or ten million contacts a year?

Open-ended or repeatable?

Building is open-ended. Most of operations is repeatable.

COUNTING HOUSE

Are you in the Counting House?

The tell: **high-volume work items** • **a legacy system of record** • **money or eligibility decisions under human approval** • **a regulator, and a waiver feed** (moratoria, hardship programs, outage credits, emergency waivers).

Airlines & travel ★

refund • rebook • waiver

P&C insurance claims ★

intake • payment • rental authorization

Health insurance payers ★

claim status • appeals • prior auth

Hospital revenue cycle

billing disputes • payment plans • assistance

Retail banking ★

card disputes • fee reversals • hardship

Credit card issuers

chargebacks • credit lines • fee waivers

Mortgage servicing ★

forbearance • loss mitigation • escrow

Consumer & student lending

deferral • payoff • repayment plans

Telecom ★

billing adjustments • outage credits • plan changes

Utilities ★

payment arrangements • disconnect holds • move-in/out

Public-sector benefits ★

eligibility • payment issues • appeals

Retail, logistics, hospitality

returns • lost-parcel claims • cancellations & credits

No contact center? It still fits: card-dispute back office, claims adjudication, underwriting ops, finance close, accounts-payable exceptions, KYC refresh, settlement breaks. ★ = strongest fit.

*In the Wheelhouse, the agent is a colleague you trust.
In the Counting House, the agent is a clerk you measure.*

Take home

- Decide which house you're building before you build it.
- Coworkers need doctrine and fences.
- Clerks need rails and statistics.
- Both need one gate and an honest record.

Want a workshop? engagements@yegge.ai

Aristocrat

Genetec

Arkana Labs

Gas City

SageOx

Tessl

Snyk

GitLab